

# Modeling time series using LSTMs and TCNs

Hajar Hajji

INRAE/IECL

Supervisors : Gabriel Destouet, Marianne Clausel, Emilie Joetzjer

- 1) Dataset
- 2) LSTM :  
Long-Short-  
Term-Memory
- 3) TCN :  
Temporal  
Convolutional  
Networks
- Conclusion

## 1) Dataset

## 2) LSTM : Long-Short-Term-Memory

## 3) TCN : Temporal Convolutional Networks

## Conclusion

- 1) Dataset
- 2) LSTM :  
Long-Short-  
Term-Memory
- 3) TCN :  
Temporal  
Convolutional  
Networks
- Conclusion

## 1) Dataset

1) Dataset

2) LSTM :  
Long-Short-  
Term-Memory3) TCN :  
Temporal  
Convolutional  
Networks

Conclusion

Variable	Description
GPP	Quantity of CO2 assimilated by the forest through photosynthesis
SWC	Soil water content
VPD	Vapor pressure deficit
Time	format : YYYY-MM-DD HH:MM:SS

Table 1: Variable Descriptions

- Mathematical formula:

$$GPP_t = \overline{GPP}_t^{365d} + \overline{GPP_t - GPP_t}^{\psi^{365d} 24h} + GPP_t$$

where

$$\overline{GPP}^T = \frac{1}{N} \sum_{k=1}^N GPP_{t+kT}$$

- Mathematical formula:

$$GPP_t = \overline{GPP}_t^{365d} + \overline{GPP_t - GPP_t}^{\psi 365d 24h} + \Delta GPP_t$$

where

$$\overline{GPP}^T = \frac{1}{N} \sum_{k=1}^N GPP_{t+kT}$$

- $\overline{GPP}_t^{365d}$  : Annual mean calculated with  $N = 365$ .
- $\overline{GPP_t - GPP_t}^{\psi 365d 24h}$  : Day-and-night cycle calculated with  $N = 365 \times 24$ .
- $\Delta GPP_t$  : Anomalies.

# GPP Decomposition

- 1) Dataset
- 2) LSTM :  
Long-Short-Term-Memory
- 3) TCN :  
Temporal  
Convolutional  
Networks
- Conclusion

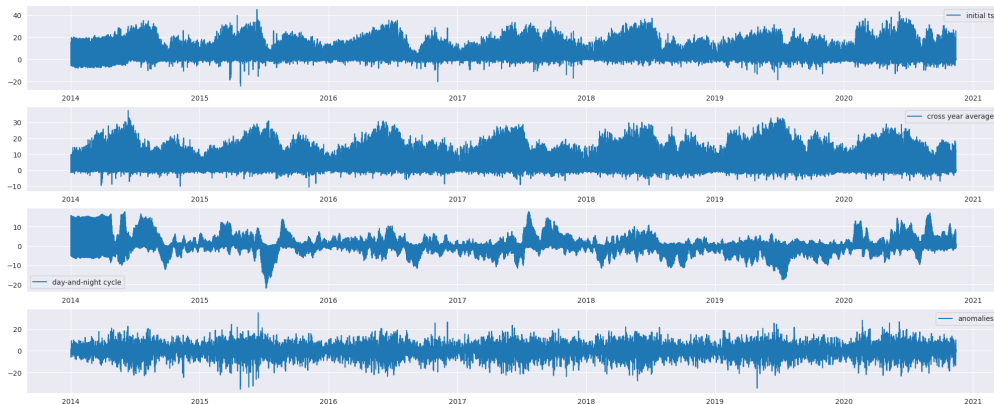


Fig. 1: GPP Decomposition

## 1) Dataset

2) LSTM :  
Long-Short-  
Term-Memory

3) TCN :  
Temporal  
Convolutional  
Networks

Conclusion




Fig. 2: Day time signal




Fig. 3: Year time signal



- 1) Dataset
- 2) LSTM :  
Long-Short-  
Term-Memory
- 3) TCN :  
Temporal  
Convolutional  
Networks
- Conclusion

## 2) LSTM : Long-Short-Term-Memory

1) Dataset

2) LSTM :  
Long-Short-  
Term-Memory

3) TCN :  
Temporal  
Convolutional  
Networks

Conclusion

- Inputs : Sequences

1) Dataset

2) LSTM :  
Long-Short-  
Term-Memory3) TCN :  
Temporal  
Convolutional  
Networks

Conclusion

- Inputs : Sequences

- Annual GPP
  - Day-and-night cycle GPP
  - Anomalies
  - Time features
- } from  $t - L$  to  $t$   
 $L$  : Sequence length

1) Dataset

2) LSTM :  
Long-Short-  
Term-Memory3) TCN :  
Temporal  
Convolutional  
Networks

Conclusion

- Inputs : Sequences

- Annual GPP
  - Day-and-night cycle GPP
  - Anomalies
  - Time features
- } from  $t - L$  to  $t$   
 $L$  : Sequence length

- Output :

1) Dataset

2) LSTM :  
Long-Short-  
Term-Memory3) TCN :  
Temporal  
Convolutional  
Networks

Conclusion

- Inputs : Sequences

- Annual GPP
  - Day-and-night cycle GPP
  - Anomalies
  - Time features
- } from  $t - L$  to  $t$   
 $L$  : Sequence length

- Output :

- Annual GPP
  - Day-and-night cycle GPP
  - Anomalies
- } at  $t + 1$

1) Dataset

2) LSTM :  
Long-Short-  
Term-Memory3) TCN :  
Temporal  
Convolutional  
Networks

Conclusion

Layer (type)	Output Shape	Param #
lstm (LSTM)	(None, 128)	69632
dense (Dense)	(None, 3)	387
Total params: 70,019		
Trainable params: 70,019		
Non-trainable params: 0		

Fig. 4: Model architecture

1) Dataset

2) LSTM :  
Long-Short-  
Term-Memory

3) TCN :  
Temporal  
Convolutional  
Networks

Conclusion

- Sequence length = 365 (7days)
- Features = 7
- Hidden units = 128
- Optimizer = Adam
- Learning rate = 0.001
- Batch size = 32

1) Dataset

2) LSTM :  
Long-Short-  
Term-Memory

3) TCN :  
Temporal  
Convolutional  
Networks

Conclusion




Fig. 5: Training loss VS Validation loss : Cross year average



1) Dataset

2) LSTM :  
Long-Short-  
Term-Memory3) TCN :  
Temporal  
Convolutional  
Networks

Conclusion




Fig. 6: Training loss VS Validation loss : Day-and-night cycle

1) Dataset

2) LSTM :  
Long-Short-  
Term-Memory3) TCN :  
Temporal  
Convolutional  
Networks

Conclusion




Fig. 7: Training loss VS Validation loss : Anomalies

1) Dataset

2) LSTM :  
Long-Short-  
Term-Memory

3) TCN :  
Temporal  
Convolutional  
Networks

Conclusion

- Sequence length = 700 (2weeks)
- Features = 7
- Hidden units = 128
- Optimizer = Adam
- Learning rate = 0.001
- Batch size = 32

1) Dataset

2) LSTM :  
Long-Short-  
Term-Memory3) TCN :  
Temporal  
Convolutional  
Networks

Conclusion




Fig. 8: Training loss VS Validation loss : Cross year average

1) Dataset

2) LSTM :  
Long-Short-  
Term-Memory

3) TCN :  
Temporal  
Convolutional  
Networks

Conclusion

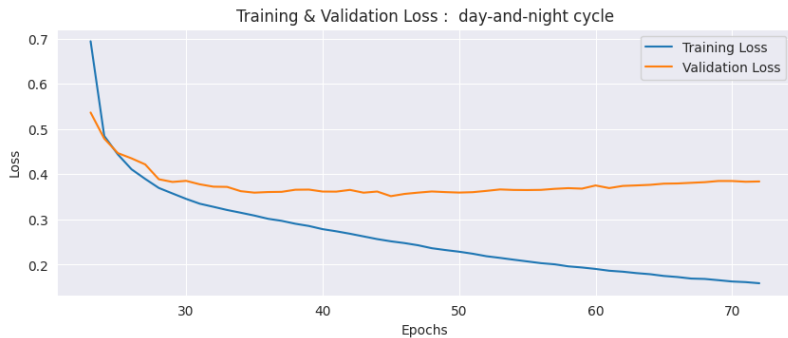


Fig. 9: Training loss VS Validation loss : Day-and-night cycle

1) Dataset

2) LSTM :  
Long-Short-  
Term-Memory3) TCN :  
Temporal  
Convolutional  
Networks

Conclusion




Fig. 10: Training loss VS Validation loss : Anomalies

1) Dataset

2) LSTM :  
Long-Short-  
Term-Memory

3) TCN :  
Temporal  
Convolutional  
Networks

Conclusion

Layer (type)	Output Shape	Param #
=====		
lstm_2 (LSTM)	(None, 365, 64)	18432
dropout_2 (Dropout)	(None, 365, 64)	0
lstm_3 (LSTM)	(None, 64)	33024
dropout_3 (Dropout)	(None, 64)	0
dense_1 (Dense)	(None, 3)	195
=====		
Total params: 51,651		
Trainable params: 51,651		
Non-trainable params: 0		

Fig. 11: Model architecture

**Learning rate = 0.0001**

**Optimizer : Adam**



**Optimizer : AdamW**





- 1) Dataset
- 2) LSTM :  
Long-Short-  
Term-Memory
- 3) TCN :  
Temporal  
Convolutional  
Networks
- Conclusion

## 3) TCN : Temporal Convolutional Networks

## Why privilege TCNs over RNNs ?

- Efficient parallelism.

## Why privilege TCNs over RNNs ?

- Efficient parallelism.
- Stable gradients (preventing vanishing/exploding gradients..)

## Why privilege TCNs over RNNs ?

- Efficient parallelism.
- Stable gradients (preventing vanishing/exploding gradients..)
- Vision across the sequence => Flexible length inputs.

The architecture of TCN consists of :

- Residual Blocks.

The architecture of TCN consists of :

- Residual Blocks.
  - Each residual block consists of two causal convolutional layers with the same dilation factor.

1) Dataset

2) LSTM :  
Long-Short-  
Term-Memory3) TCN :  
Temporal  
Convolutional  
Networks

Conclusion

The architecture of TCN consists of :

- Residual Blocks.
  - Each residual block consists of two causal convolutional layers with the same dilation factor.
  - Introduce complexity to each layer (Activation, Normalization, Regularization)

The architecture of TCN consists of :

- Residual Blocks.
  - Each residual block consists of two causal convolutional layers with the same dilation factor.
  - Introduce complexity to each layer (Activation, Normalization, Regularization)
  - The output of the convolutional layers is added to the input of the residual block through a residual connection.




The architecture of TCN consists of :

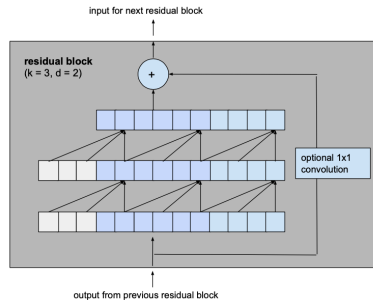
- Residual Blocks.
  - Each residual block consists of two causal convolutional layers with the same dilation factor.
  - Introduce complexity to each layer (Activation, Normalization, Regularization)
  - The output of the convolutional layers is added to the input of the residual block through a residual connection.
- Sequential arrangement of residual blocks with increasingly dilations (with blocks).

# Residual blocks

- 1) Dataset
- 2) LSTM :  
Long-Short-Term-Memory
- 3) TCN :  
Temporal Convolutional Networks
- Conclusion



(a) Causal convolutional layer



(b) Residual block

Fig. 13: Residual blocks

- 1) Dataset
- 2) LSTM :  
Long-Short-Term-Memory
- 3) TCN :  
Temporal  
Convolutional  
Networks
- Conclusion

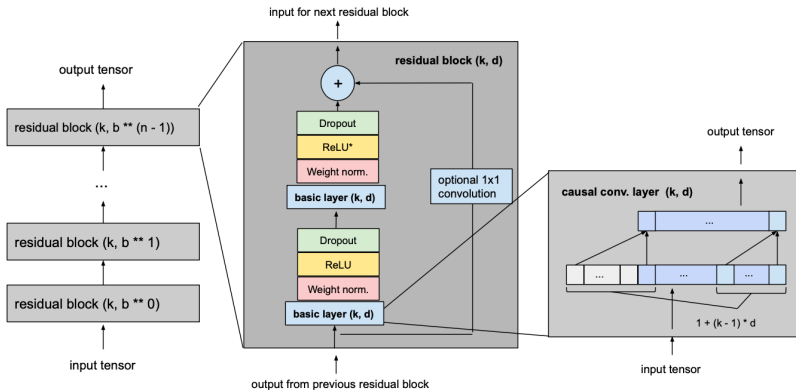


Fig. 14: The entire model

- Remove the first year (2014) and last year (2022) (due to too many SWC missing values).



Fig. 15: Missing SWC values on September 2016

- Remove the first year (2014) and last year (2022) (due to too many SWC missing values).
- Filling SWC gaps using interpolation and sometimes backward filling.



Fig. 15: Missing SWC values on September 2016

1) Dataset

2) LSTM :  
Long-Short-  
Term-Memory3) TCN :  
Temporal  
Convolutional  
Networks

Conclusion

- Remove the first year (2014) and last year (2022) (due to too many SWC missing values).
- Filling SWC gaps using interpolation and sometimes backward filling.
- Penalize the loss function based on the number of SWC values that have been filled.



Fig. 15: Missing SWC values on September 2016

- Inputs : Sequences

- Inputs : Sequences

- GPP
  - VPD
  - SWC
  - Time features
- } from  $t - L$  to  $t$   
 $L$  : Sequence length



- Inputs : Sequences

- GPP	}	from $t - L$ to $t$ $L$ : Sequence length
- VPD		
- SWC		
- Time features		

- Output : GPP at  $t+1$

1) Dataset

2) LSTM :  
Long-Short-  
Term-Memory

3) TCN :  
Temporal  
Convolutional  
Networks

Conclusion

- Sequence length = 400 (>one year)
- Features = 7
- Channels = [32, 64, 128]
- Kernel size = 3
- Learning rate = 0.001
- Dropout = 0.2
- Batch size = 32
- Loss function : MSE (with penalization)

- 1) Dataset
- 2) LSTM :  
Long-Short-  
Term-Memory
- 3) TCN :  
Temporal  
Convolutional  
Networks
- Conclusion




Fig. 16: Training loss VS Validation loss using TCN

- 1) Dataset
- 2) LSTM :  
Long-Short-  
Term-Memory
- 3) TCN :  
Temporal  
Convolutional  
Networks

Conclusion

## Conclusion

Overall good performance of TCN.

Perks :

- Fast & optimal convergence.
- Not time-consuming.
- No need for manual GPP decomposition.
- Interpretability.

Drawbacks :

- Requires long sequences to capture temporal patterns  
=> significant computational resources.